

Lecture 20: Introduction to Deep Learning

- Basics of feedforward Neural Networks
- Back-Propagation
- Convolutional Neural Net
- Some design examples

Lecture notes on Deep Learning

VCLA@UCLA


Logistic Regression and Logistic Function

Logistic regression has a learned parameter vector W . On input X , label $Y \in \{-1, 1\}$, the logistic regression is defined as generalized linear model with link function:

$$\log \frac{P(Y = 1|X)}{P(Y = -1|X)} = WX$$

$$P(Y = 1|X) = \frac{1}{1 + \exp(-WX)}$$

In early days, neural networks use logistic function as activation function:



Lecture notes on Deep Learning

VCLA@UCLA


Neural Network

Example: 3 layer network with classification:

In multi-class classification, label $Y \in \{1, 2, \dots, n\}$ has n classes. In modern deep learning, softmax layer as multinomial logistic regression is used to compute the probability of each class, and output $Y = \underset{Y}{\operatorname{argmax}} P(Y|X)$

$$P(Y = i|X) = \frac{\exp(X_i^{(3)})}{\sum_{j=1}^n \exp(X_j^{(3)})}, i = 1, 2, \dots, n$$


Where $X_i^{(3)}$ is the output for i -th class in the third layer.



Lecture notes on Deep Learning

VCLA@UCLA

Training a Neural Network



Given training set $(X_1, Y_1), (X_2, Y_2), \dots, (X_M, Y_M), \dots$, adjust parameters W (for every node) to maximize likelihood:

$$P(Y = i|X; W) = \frac{\exp(X_i^{(L)})}{\sum_{j=1}^n \exp(X_j^{(L)})}$$

$X^{(L)} = \langle W, X^{(L-1)} \rangle$: final layer.

$X^{(l)}$: output from l -th layer, $l=1, 2, \dots, L-1$.

$X^{(0)} = X$ is the input data.

Use gradient ascent. "Back-propagation" algorithm.

In 1980's, the reconstruction error is used as loss function $\sum_{m=1}^M (Y_m - f(X_m))^2$

Lecture notes on Deep Learning

VCLA@UCLA

Back-propagation

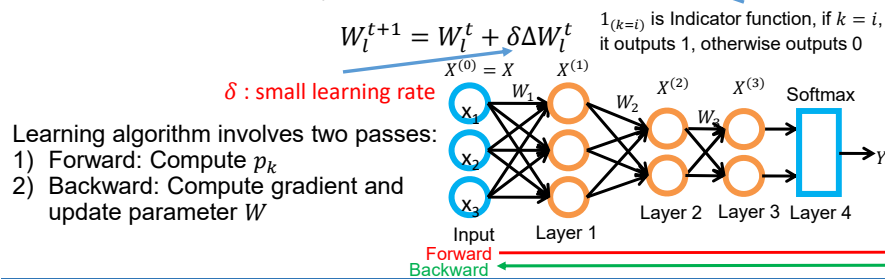
Let $X^{(l)}$ denote the l -th layer output, $l = 1, \dots, L$

The gradient for l -th layer parameter:

$$\Delta W_l^t = \frac{\partial P(Y=i|X;W^t)}{\partial W_l^t} = \frac{\partial P(Y=i|X;W^t)}{\partial X^{(L)}} \cdot \frac{\partial X^{(L)}}{\partial X^{(L-1)}} \cdots \frac{\partial X^{(L+1)}}{\partial W_l^t}$$

$$\frac{\partial P(Y=i|X;W^t)}{\partial X^{(L)}} = [p_1 \cdot (1_{(1=i)} - p_i), \dots, p_k \cdot (1_{(k=i)} - p_i), \dots, p_n \cdot (1_{(n=i)} - p_i)]$$

$$p_k = \frac{\exp(X_k^{(L)})}{\sum_{j=1}^n \exp(X_j^{(L)})}, k = 1, 2, \dots, n$$



Lecture notes on Deep Learning

VCLA@UCLA

Back-propagation

Chain rule part:

Let $X^{(l)}$ denote the l -th layer output, $l = 1, \dots, L$

$$X^{(l)} = f_l(X^{(l-1)}; W_l) = \frac{1}{1 + \exp(-W_l X^{(l-1)})}$$

$$\frac{\partial X^{(l)}}{\partial X^{(l-1)}} = \frac{1}{1 + \exp(-W_l X^{(l-1)})} \frac{\exp(-W_l X^{(l-1)})}{1 + \exp(-W_l X^{(l-1)})} W_l$$

$$\frac{\partial X^{(l)}}{\partial W_l} = \frac{1}{1 + \exp(-W_l X^{(l-1)})} \frac{\exp(-W_l X^{(l-1)})}{1 + \exp(-W_l X^{(l-1)})} X^{(l-1)}$$

The gradient in deep learning is unstable, tending to either explode (large gradient) or vanish (small gradient) in early layers. Batch-normalization is used to normalize output of each layer with a mean of 0 and a standard deviation of 1.

Rumelhart, David E.; Hinton, Geoffrey E.; Williams, Ronald J. (8 October 1986). "Learning representations by back-propagating errors". *Nature*. **323** (6088): 533–536.

Lecture notes on Deep Learning


VCLA@UCLA

Convolutional Neural Network

The replicated feature approach.

- Use many different copies of the same feature detector with different positions.
 - Could also replicate across scale and orientation (tricky and expensive)
 - Replication greatly reduces the number of free parameters to be learned.
- Use several different feature types, each with its own map of replicated detectors.
 - Allows each patch of image to be represented in several ways.

The red connections all have the same weight with constraint.




LeCun, Yann; Léon Bottou; Yoshua Bengio; Patrick Haffner (1998). "Gradient-based learning applied to document recognition". *Proceedings of the IEEE*. 86 (11): 2278–2324. doi:10.1109/5.726791.

Lecture notes on Deep Learning


VCLA@UCLA

Convolution



Let F denote the filter, the convolution is defined:

$$[F * X](y) = \sum_{x \in S} W_x X(y + x)$$




The convolution layer's parameters consist of a set of learnable filters. Every filter is small spatially (along width and height), but extends through the full depth of the input volume.

Lecture notes on Deep Learning

VCLA@UCLA

Nonlinear Transformation

Rectified linear units (ReLU):
the rectifier is an activation function defined as:



$$f(X) = \max(X, 0)$$


Motivated by strong biological motivations and mathematical justifications

It has been widely used and the most popular activation function in ConvNet.

R Hahnloser, R. Sarpeshkar, M A Mahowald, R. J. Douglas, H.S. Seung (2000). *Digital selection and analogue amplification coexist in a cortex-inspired silicon circuit*. *Nature*. **405**. pp. 947–951.

Max pooling:

Max pooling is a form of non-linear down-sampling. It partitions the input image into a set of non-overlapping rectangles and, for each such sub-region, outputs the maximum.

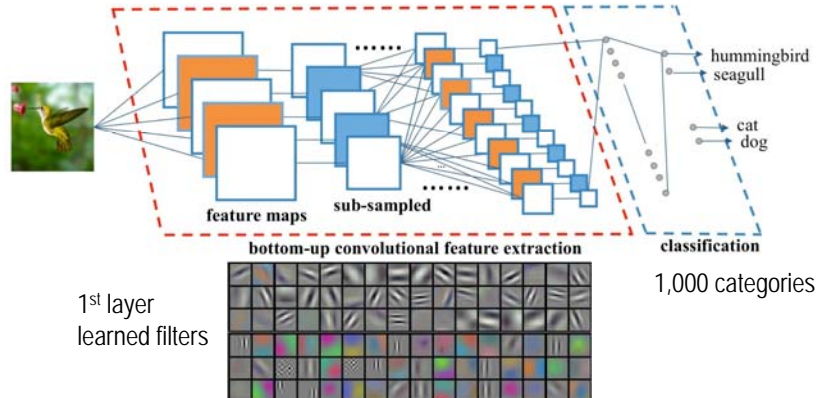


Lecture notes on Deep Learning

VCLA@UCLA

Convolutional Neural Network

Example of modern ConvNet and learned 1st layer filters:



Krizhevsky, A.; Sutskever, I.; Hinton, G. E. (2012). "Imagenet classification with deep convolutional neural networks". *Advances in Neural Information Processing Systems*. 1: 1097–1105.

Lecture notes on Deep Learning

VCLA@UCLA

Convolutional Neural Network

Formally, a Convolutional Neural Network (ConvNet) is used to define a mapping:

$$Y = f(X; W)$$

where W denotes the parameters. $f(\cdot)$ is a composition of L layers of linear transformations followed by element-wise non-linear transformations:

$$X^{(l)} = f_l(W_l X^{(l-1)})$$

where $l = 1, \dots, L$, $X^{(0)} = X$, $X^{(L)} = Y$, W_l is the weights matrix. f_l is element-wise nonlinear transformation.

Key factors why modern ConvNets work (empiric observations)

- ReLU activation function.
- Multinomial logistic regression
- Fully connected layer
- Dropout (Remove 50% of connection)


Hinton, Geoffrey E.; Srivastava, Nitish; Krizhevsky, Alex; Sutskever, Ilya; Salakhutdinov, Ruslan R "Dropout: A Simple Way to Prevent Neural Networks from Overfitting". Jmlr.org. Retrieved July 26, 2015.

Lecture notes on Deep Learning

VCLA@UCLA

What does replicating the feature detectors achieve?

- **Equivariant activities:** Replicated features do **not** make the neural activities invariant to translation. The activities are equivariant.



- **Invariant knowledge:** If a feature is useful in some locations during training, detectors for that feature will be available in all locations during testing.

Lecture notes on Deep Learning

VCLA@UCLA

Pooling the outputs of replicated feature detectors

- Get a small amount of translational invariance at each level by averaging four neighboring replicated detectors to give a single output to the next level.
 - This reduces the number of inputs to the next layer of feature extraction, thus allowing us to have many more different feature maps.
 - Taking the maximum of the four works slightly better.
- **Problem:** After several levels of pooling, we have lost information about the precise positions of things.
 - This makes it impossible to use the precise spatial relationships between high-level parts for recognition.

Lecture notes on Deep Learning

VCLA@UCLA

LeNet

- Yann LeCun and his collaborators developed a really good recognizer for handwritten digits by using backpropagation in a feedforward net with:
 - Many hidden layers
 - Many maps of replicated units in each layer.
 - Pooling of the outputs of nearby replicated units.
 - A wide net that can cope with several characters at once even if they overlap.
 - A clever way of training a complete system, not just a recognizer.
- This net was used for reading ~10% of the checks in North America.

LeCun, Yann; Léon Bottou; Yoshua Bengio; Patrick Haffner (1998). "Gradient-based learning applied to document recognition". *Proceedings of the IEEE*. 86 (11): 2278–2324. doi:10.1109/5.726791.

Lecture notes on Deep Learning

VCLA@UCLA

Priors and Prejudice

- We can put our prior knowledge about the task into the network by designing appropriate:
 - Connectivity.
 - Weight constraints.
 - Neuron activation functions
- This is less intrusive than hand-designing the features.
 - But it still prejudices the network towards the particular way of solving the problem that we had in mind.
- Alternatively, we can use our prior knowledge to create a whole lot more training data.
 - This may require a lot of work (Hofman&Tresp, 1993)
 - It may make learning take much longer.
- It allows optimization to discover clever ways of using the multi-layer network that we did not think of.
 - And we may never fully understand how it does it.

Hinton, Geoffrey. Notes for ConvNet.

Lecture notes on Deep Learning


VCLA@UCLA

AlexNet

params	AlexNet	FLOPs
4M	FC 1000	4M
16M	FC 4096 / ReLU	16M
37M	FC 4096 / ReLU	37M
	Max Pool 3x3s2	
442K	Conv 3x3s1, 256 / ReLU	74M
1.3M	Conv 3x3s1, 384 / ReLU	112M
884K	Conv 3x3s1, 384 / ReLU	149M
	Max Pool 3x3s2	
	Local Response Norm	
307K	Conv 5x5s1, 256 / ReLU	223M
	Max Pool 3x3s2	
	Local Response Norm	
35K	Conv 11x11s4, 96 / ReLU	105M

Model	Top-1 (val)	Top-5 (val)	Top-5 (test)
<i>SIFT + FVs [7]</i>	—	—	26.2%
1 CNN	40.7%	18.2%	—
5 CNNs	38.1%	16.4%	16.4%
1 CNN*	39.0%	16.6%	—
7 CNNs*	36.7%	15.4%	15.3%

Table 2: Comparison of error rates on ILSVRC-2012 validation and test sets. In *italics* are best results achieved by others. Models with an asterisk* were "pre-trained" to classify the entire ImageNet 2011 Fall release. See Section 6 for details.



Krizhevsky, A.; Sutskever, I.; Hinton, G. E. (2012). "Imagenet classification with deep convolutional neural networks". *Advances in Neural Information Processing Systems*. 1: 1097–1105.

Lecture notes on Deep Learning

VCLA@UCLA

